

Final Project – Text Generation

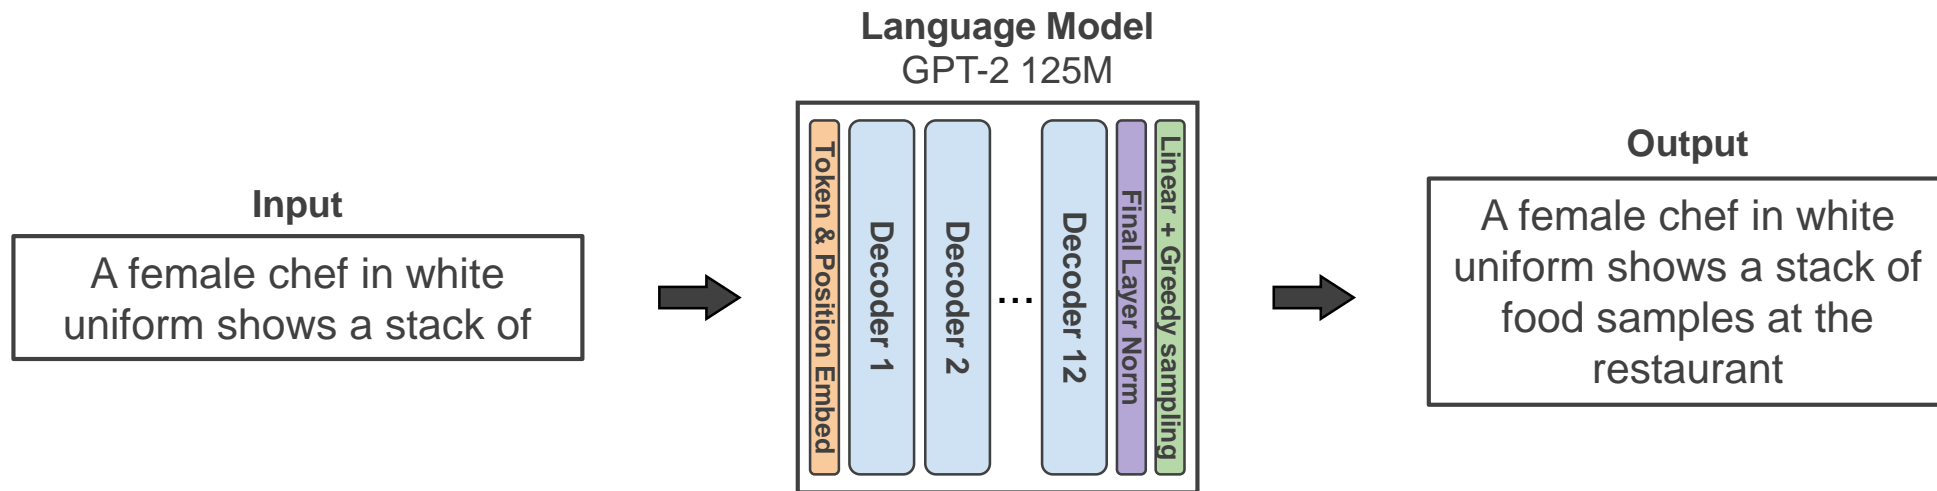
Project due: 2024. 06. 17. 11:59 PM

Last updated: 2024. 06. 10. 21:00 PM

Project Goal

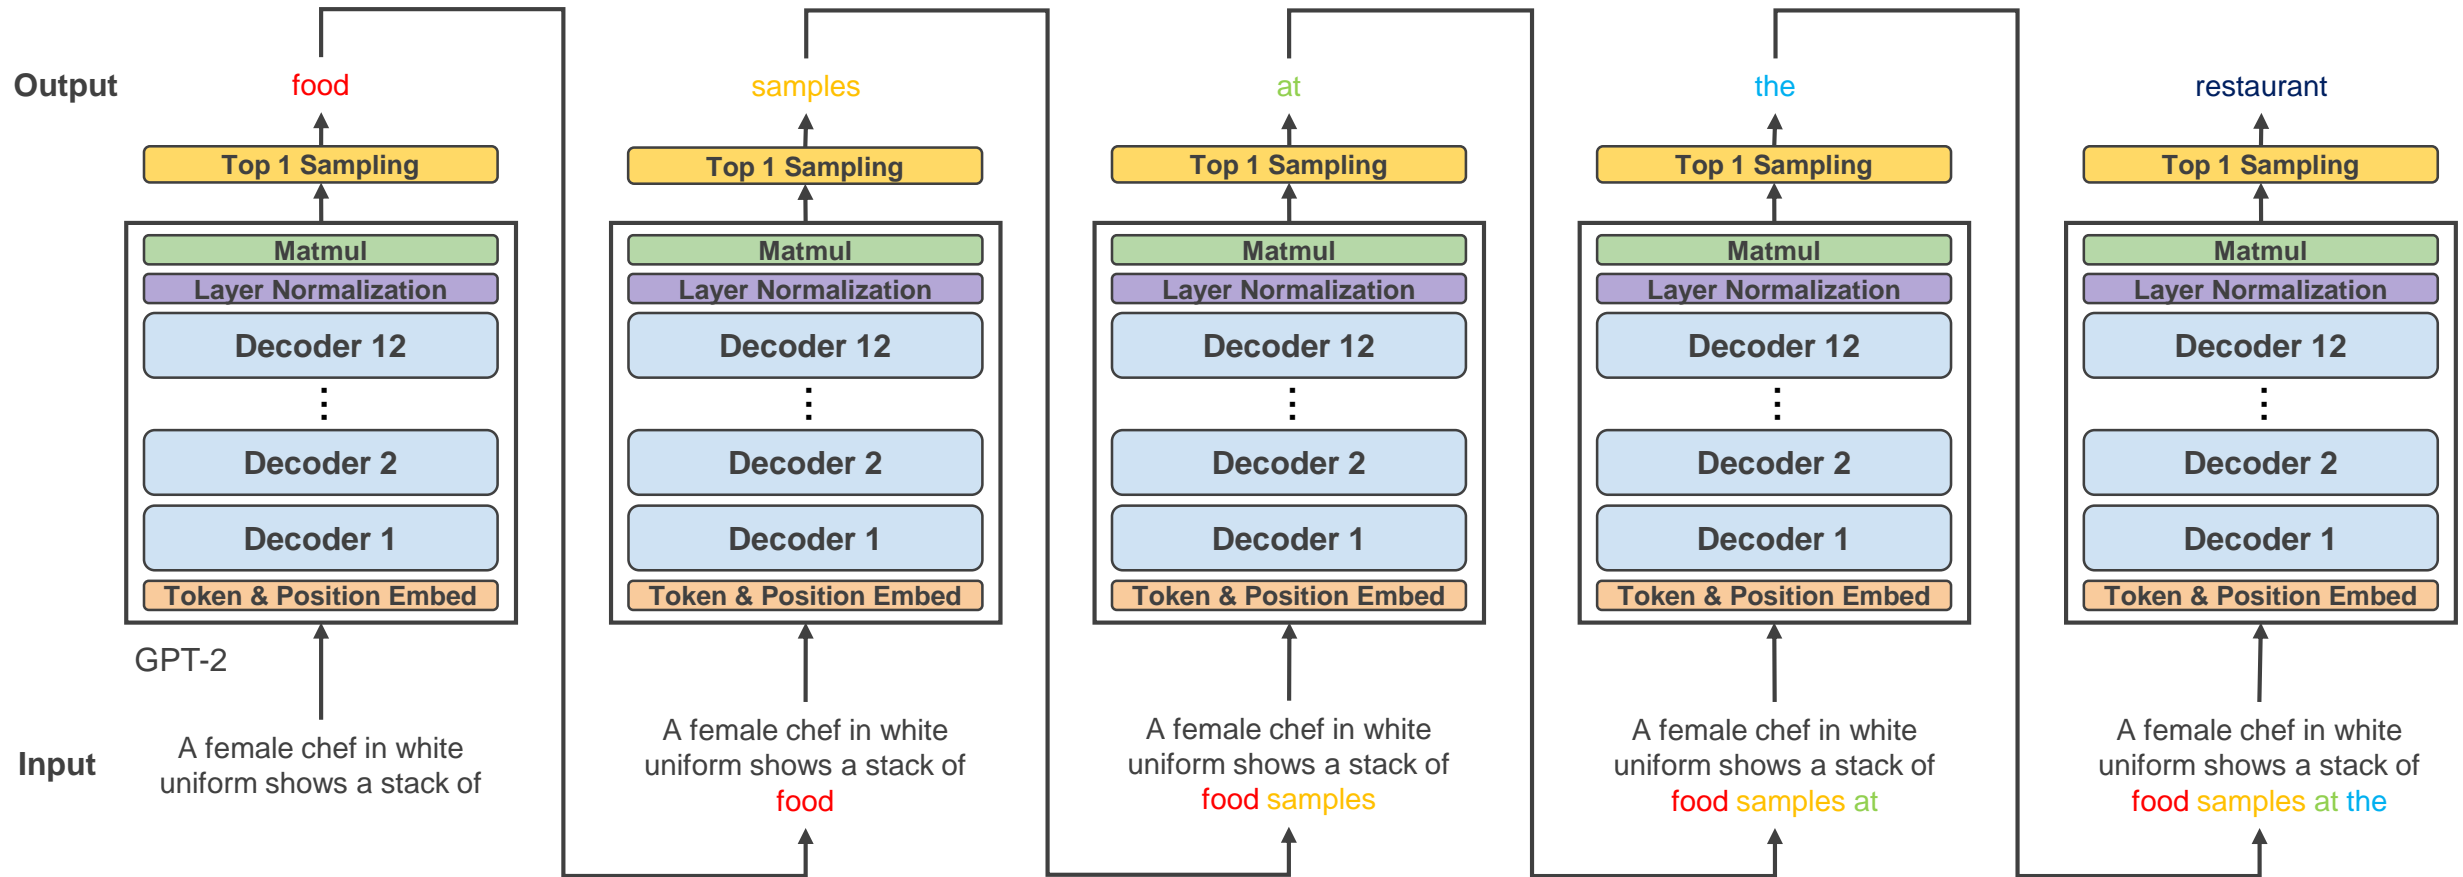
Optimize GPT-2 125M text generation

- You are given a sequential (CPU, single thread) code
- Parallelize and optimize the code across 4 nodes (16 GPUs in total)



Background and Skeleton Code

Text Generation using GPT-2



Skeleton Code

/shpc24/skeleton/final-project/

include/

tensor.h

Tensor structure definition

layer.h

Layer definitions

model.h

GPT-2 model configuration and interfaces, **do not modify**

src/

tensor.cu

Tensor structure implementation

layer.cu

DNN layer implementation, where actual computations are done

model.cu

GPT-2 model implementation using DNN layers

main.cpp

Driver code, **do not modify**

data/

Directory to store input/output

Makefile

Makefile, **do not modify**

run.sh

How to Run (1)

Compile with make command

```
shpcta@login0:~/final-project$ make -j
mkdir -p obj
mpic++ -std=c++14 -O3 -Wall -march=native -mavx2 -fopenmp -I/usr/local/cuda/
/usr/local/cuda/bin/nvcc -Xcompiler=-std=c++14 -Xcompiler=-O3 -Xcompiler
iler=-Iinclude -c -o obj/model.o src/model.cu
/usr/local/cuda/bin/nvcc -Xcompiler=-std=c++14 -Xcompiler=-O3 -Xcompiler
iler=-Iinclude -c -o obj/tensor.o src/tensor.cu
/usr/local/cuda/bin/nvcc -Xcompiler=-std=c++14 -Xcompiler=-O3 -Xcompiler
iler=-Iinclude -c -o obj/layer.o src/layer.cu
cc -std=c++14 -O3 -Wall -march=native -mavx2 -fopenmp -I/usr/local/cuda/inc
art -lm -lmpi -lmpi_cxx
shpcta@login0:~/final-project$
```

How to Run (2)

Run with run.sh script

- Make sure to understand the script
- Check out the program options

```
shpcta@login0:~/chundoong-lab-ta/SHPC2024/final-project$ ./run.sh -h
salloc: Pending job allocation 527267
salloc: job 527267 queued and waiting for resources
salloc: job 527267 has been allocated resources
salloc: Granted job allocation 527267
Usage: ./main [-i 'pth'] [-p 'pth'] [-o 'pth'] [-a 'pth'] [-t 'tokens'] [-n 'prompts'] [-v] [-h]
Options:
  -i: Input binary path (default: data/input.bin)
  -p: Model parameter path (default: /shpc24/project_model_parameters.bin)
  -o: Output binary path (default: output.bin)
  -a: Answer binary path (default: data/answer.bin)
  -n: Number of input prompts (default: 1)
  -t: Number of tokens to generate (default: 8)
  -v: Enable validation (default: OFF)
  -h: Print manual and options (default: OFF)
salloc: Relinquishing job allocation 527267
shpcta@login0:~/chundoong-lab-ta/SHPC2024/final-project$
```

How to Run (3)

Run example

- 4 prompts
- Generate 8 tokens per prompt
- Validate output

```
shpcta@login0:~/final-project$ ./run.sh -v -n 4 -t 8
salloc: Pending job allocation 527342
salloc: job 527342 queued and waiting for resources
salloc: job 527342 has been allocated resources
salloc: Granted job allocation 527342

=====
Model: GPT-2 125M
=====

Validation: ON
Number of Prompts: 4
Number of Tokens to generate: 8
Input binary path: ./data/input.bin
Model parameter path: /shpc24/project_model_parameters.bin
Answer binary path: ./data/answer.bin
Output binary path: ./data/output.bin
=====

Initializing input and parameters...Done
Generating tokens...Done!
Elapsed time: 29.984716 (sec)
Throughput: 1.067210 (tokens/sec)
Finalizing...Done
Saving output to ./data/output.bin...Done
Validation...PASS
salloc: Relinquishing job allocation 527342
```


How to Run (4)

- Skeleton code uses tokenized texts (integers)
- Use shpc24-project-bin2text utility to see the actual text

```
shpcta@login0:~/final-project$ shpc24-bin2text
Usage: shpc24-bin2text <input_path> <output_path> <prompt_count (optional)>
e.g., shpc24-bin2text ./data/input.bin ./data/output.bin 1
<input_path>: Path to the input binary file
<output_path>: Path to the output binary file
<prompt_count>: Number of prompts to display (optional)
shpcta@login0:~/final-project$ shpc24-bin2text ./data/input.bin ./data/output.bin 4
Loading GPT-2 tokenizer...
Prompt # 1
Input Prompt: Then, the man writes over the snow covering the window of a car, and
Generated Output: the woman, who is sitting on the

Prompt # 2
Input Prompt: A female chef in white uniform shows a stack of baking pans in a large kitchen
Generated Output: . She is wearing a white dress with

Prompt # 3
Input Prompt: A female chef in white uniform shows a stack of baking pans in a large kitchen
Generated Output: . She is wearing a white dress with


Prompt # 4
Input Prompt: A tray of potatoes is loaded into the oven and removed. A large tray of
Generated Output: potatoes is placed on the stovetop and
```

Tensor

Data structure used in deep learning


- High-dimensional matrix
- Need to understand how the actual data is stored
- Defined in `include/tensor.h`, implemented in `src/tensor.cu`

1D Tensor
(e.g., Vector)




Shape = {4}

2D Tensor
(e.g., Matrix)




Shape = {2, 4}

3D Tensor
(e.g., RGB image)



Shape = {3, 2, 4}

4D Tensor
(e.g., CNN filter)



Shape = {2, 3, 2, 4}

Layers

Unit operations where the actual computations are done

- We use these layers to construct GPT-2 model
- Layers in GPT-2 are classified into four types
 1. Matrix multiplication
 2. Data movement
 3. Elementwise
 4. Other
- Defined in `include/layer.h`, implemented in `src/layer.cu`

Layers (cont'd)

Type 1: Matrix multiplication

1. Linear
 - `void linear(in, w, b, out)`
2. Matmul
 - `void matmul(in1, in2, out)`

Type 2: Data movement operations

1. Copy
 - `void copy(in, out)`
2. Transpose
 - `void transpose(in, out)`
3. Split QKV
 - `void split_qkv(in, out)`
4. Split Attention Heads
 - `void split_head(in, num_heads, out)`
5. Concat Attention Heads
 - `void concat_head(in, out)`
6. Extract QKV
 - `void extract_qkv(in, head_idx, num_heads, q, k, v)`
7. Merge Attention Heads
 - `void merge_head(in, head_idx, num_heads, out)`
8. Token Embeddig & Positional Embedding
 - `void token_pos_embedding(in, wte, wpe, out)`

Layers (cont'd)

Type 3: Elementwise Operations

1. GELU - `void gelu(inout)`
2. Add - `void add(inout, x)`
3. Scaling - `void scaling(inout, scale)`

Type 4: Other Operations

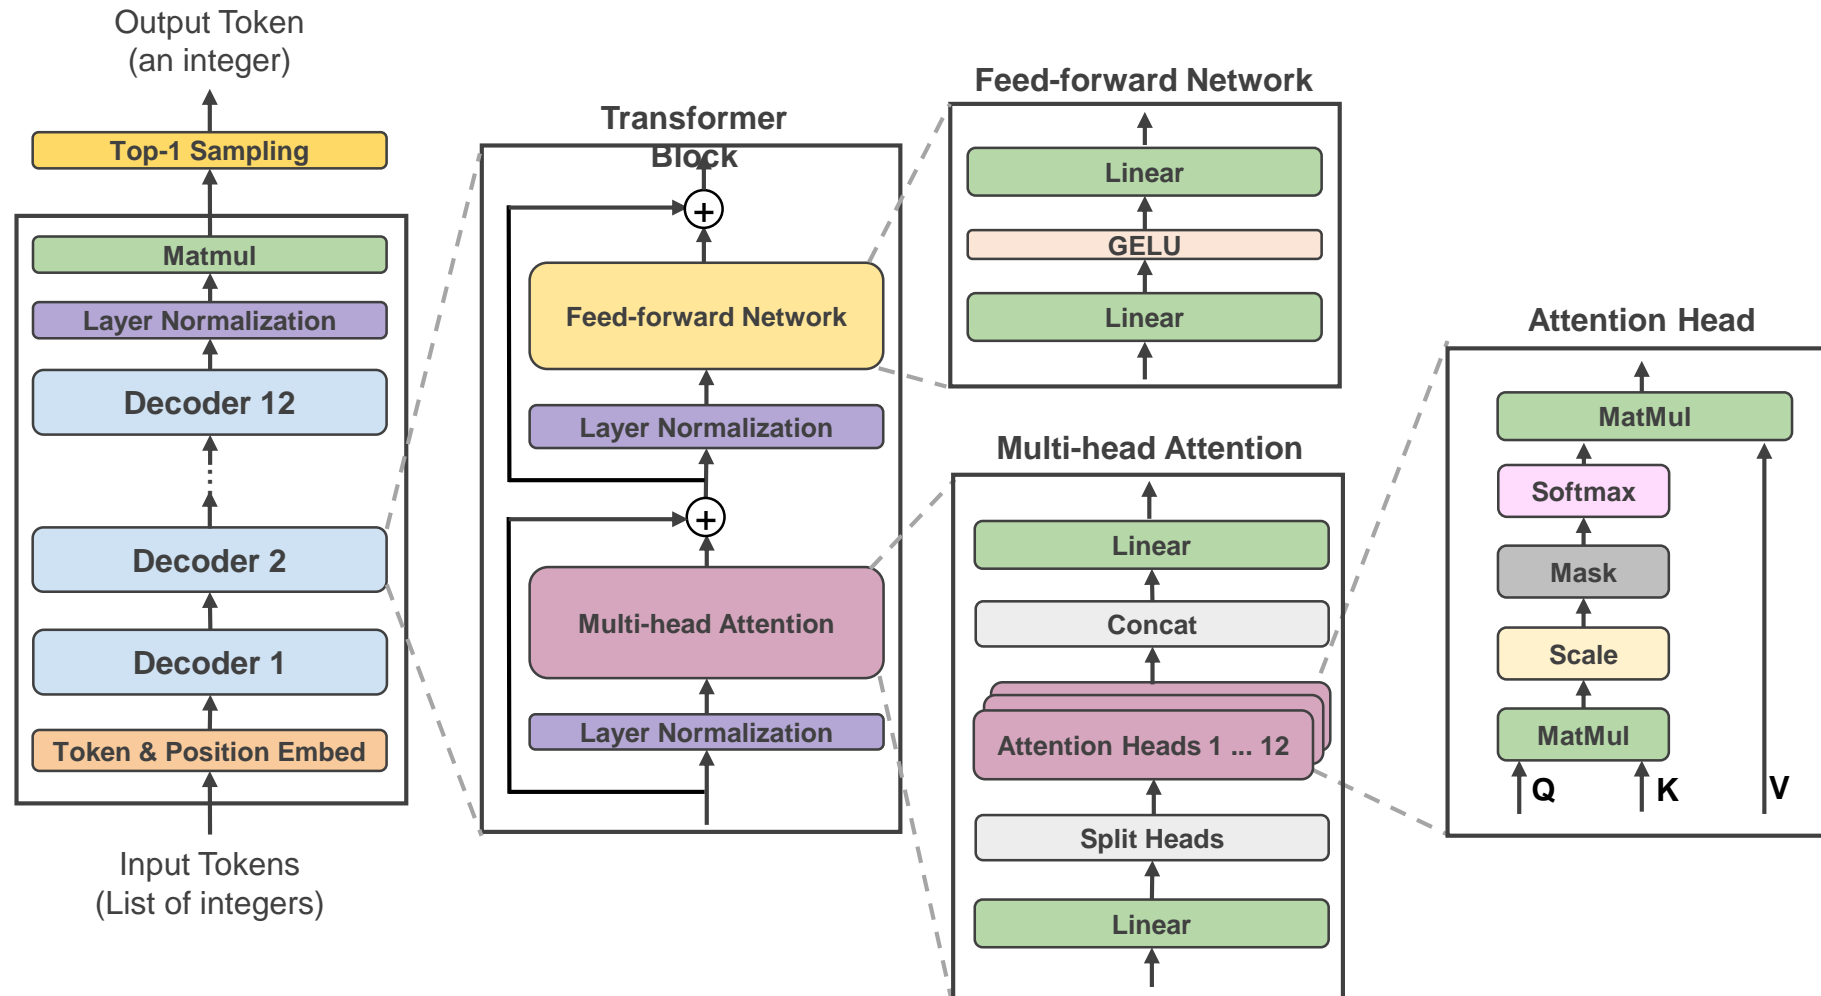
1. Softmax - `void softmax(inout)`
2. Layer Normalization - `void layer_norm(inout, gamma, beta)`
3. Mask Generation - `void generate_mask(inout)`
4. Top 1 Sampling - `int top1_sampling(in)`

Model

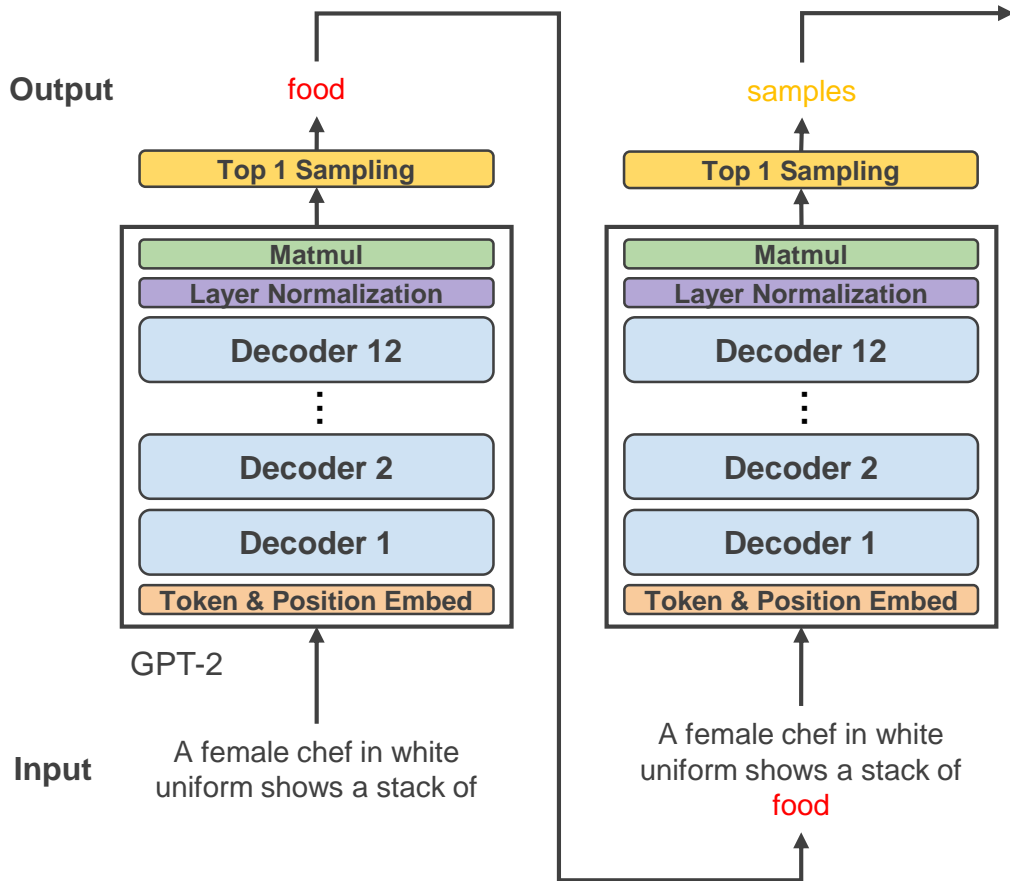
GPT-2 model architecture and interface are defined in `include/model.h`, implemented in `src/model.cu`

- `void alloc_and_set_parameters(float *param)`
 - Initialization of the model
 - Allocate memory for the model parameters, read the parameter file, and set
- `void generate_tokens(int *input, int *output, size_t n_prompt, size_t n_token)`
 - Main body of text generation using GPT-2
 - Our optimization target
 - Takes `n_prompt` prompts of the length 16, generate 8 tokens for each prompt
- `void free_parameters()`
 - Finalize the model
 - Free up the allocated memory for the model parameters

GPT-2 Model Architecture Breakdown



Model Architecture Breakdown – Text Generation Loop



```
void generate_tokens(int *input, int *output, size_t n_prompt, size_t n_token) {
    int mpi_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &mpi_rank);
    if (mpi_rank == 0) {
        int prompt_size = input_tokens;

        for (size_t p = 0; p < n_prompt; p++) {
            prompt_size = input_tokens;

            vector<int> input_prompt(prompt_size);
            memcpy(input_prompt.data(), input + p * prompt_size,
                   prompt_size * sizeof(int));

            for (size_t t = 0; t < n_token; t++) {
                alloc_activations(prompt_size);

                token_pos_embedding(input_prompt, wte, wpe, embd_a);

                for (size_t l = 0; l < NUM_LAYER; l++) {
                    transformer_block(embd_a, attn_b[l], attn_w[l], proj_b[l], proj_w[l],
                                     ln_1_b[l], ln_1_g[l], ln_2_b[l], ln_2_g[l],
                                     mlp1_b[l], mlp1_w[l], mlp2_b[l], mlp2_w[l],
                                     transformer_block_a);

                    copy(transformer_block_a, embd_a);
                }

                layer_norm(embd_a, ln_f_g, ln_f_b);

                transpose(wte, wte_transposed_a);
                matmul(embd_a, wte_transposed_a, logit_a);

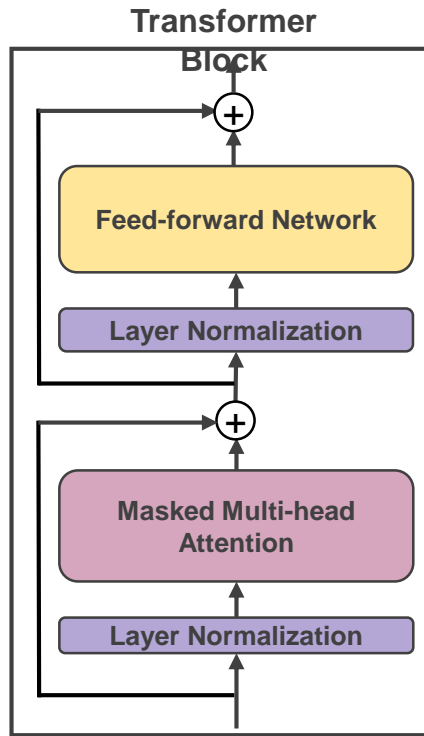
                int next_token_id = top1_sampling(logit_a);

                input_prompt.push_back(next_token_id);
                prompt_size += 1;

                output[p * n_token + t] = next_token_id;

                free_activations();
            }
        }
    }
}
```


Model Architecture Breakdown – Transformer Block



```
void transformer_block(Activation *in, Parameter *attn_b, Parameter *attn_w,
                      Parameter *proj_b, Parameter *proj_w, Parameter *ln_1_b,
                      Parameter *ln_1_g, Parameter *ln_2_b, Parameter *ln_2_g,
                      Parameter *mlp1_b, Parameter *mlp1_w, Parameter *mlp2_b,
                      Parameter *mlp2_w, Activation *out){
    copy(in, residual_a);

    layer_norm(in, ln_1_g, ln_1_b);

    mha(in, attn_b, attn_w, proj_b, proj_w, mha_out_a);

    add(mha_out_a, residual_a);


    copy(mha_out_a, residual_a);

    layer_norm(mha_out_a, ln_2_g, ln_2_b);

    ffn(mha_out_a, mlp1_w, mlp1_b, mlp2_w, mlp2_b, out);

    add(out, residual_a);
}
```


Model Architecture Breakdown – Feed-forward Network



```
void ffn(Activation *in, Parameter *mlp1_w, Parameter *mlp1_b,  
         Parameter *mlp2_w, Parameter *mlp2_b, Activation *out) {  
    linear(in, mlp1_w, mlp1_b, ffn_proj_a);  
  
    gelu(ffn_proj_a);  
  
    linear(ffn_proj_a, mlp2_w, mlp2_b, out);  
}
```

Model Architecture Breakdown – MHA and Attention Head

Multi-Head Attention



```
void mha(Activation *in, Parameter *attn_b, Parameter *attn_w,
        Parameter *proj_b, Parameter *proj_w, Activation *out) {
    linear(in, attn_w, attn_b, mha_qkv_proj_a);

    split_qkv(mha_qkv_proj_a, mha_split_qkv_a);

    split_head(mha_split_qkv_a, NUM_HEAD, mha_split_head_a);

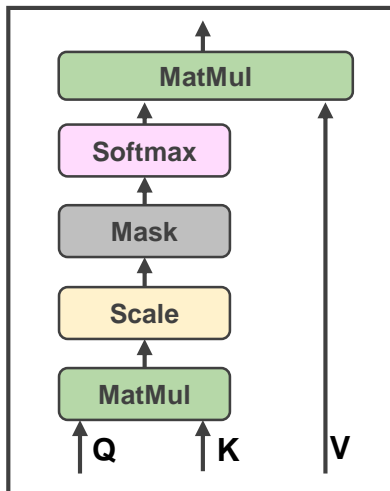
    generate_mask(mha_mask_a);

    for (size_t idx = 0; idx < NUM_HEAD; idx++) {
        extract_qkv(mha_split_head_a, idx, NUM_HEAD, mha_q_a, mha_k_a, mha_v_a);
        attention(mha_q_a, mha_k_a, mha_v_a, mha_mask_a, mha_attn_out_a);
        merge_head(mha_attn_out_a, idx, NUM_HEAD, mha_merge_head_a);
    }

    concat_head(mha_merge_head_a, mha_concat_head_a);

    linear(mha_concat_head_a, proj_w, proj_b, out);
}
```

Attention Head



```
void attention(Activation *q, Activation *k, Activation *v, Activation *mask,
              Activation *out) {
    transpose(k, k_transposed_a);
    matmul(q, k_transposed_a, attn_score_a);

    scaling(attn_score_a, (1.0 / sqrt(k->shape[1])));

    add(attn_score_a, mask);

    softmax(attn_score_a);

    matmul(attn_score_a, v, out);
}
```

Project Grading

Grading

- **Performance (80%)** - Program throughput (generated tokens/sec)
 - Prompt (input sequence) length is fixed to 16 tokens. Generate 8 tokens for each prompt
 - You can choose the number of prompts. Maximum allowed is 8192
 - 1st to 4th place get full points, 10%p deduction for each time your code is 2x slower than 4th place
 - Zero point if validation fails
- **Report (20%)**
 - Filename is report.pdf
 - Less than 5 pages, no restriction in format
 - Write concisely about your parallelization & optimization methods
 - **Make sure to include how to run your program and the expected performance result**

Restrictions

- Only the libraries we've covered in the class are allowed
 - x86 intrinsics, Pthread, OpenMP, OpenMPI, CUDA
- Any other external libraries are not allowed
 - cuBLAS, cuDNN, MAGMA, BLIS, PyTorch, Intel MKL, Tensorflow ...
- Any modification that changes the program logic is not allowed
 - You cannot
 - Do any computation in initialization phase
 - Use different model or text generation algorithm that makes the same output
 - ~~• Skip some operations that does not affect the final output~~
 - You can
 - Change memory layout, change loop order, pad data, add some auxiliary operations, kernel fusion, ...
- Use eTL board if you have any questions

Submission

- Deadline: 2024. 06. 17. 11:59 PM
 - You cannot use grace days
 - This is the strict deadline, we can't extend it anymore
- Submit the following files using `shpc-submit`
 - `tensor.h`, `tensor.cu`, `layer.h`, `layer.cu`, `model.cu`, `run.sh`
 - `report.pdf`
 - The other files will be overwritten by the skeleton code
- Start early!
 - The cluster server may be (will be) overloaded near the deadline
 - We can't extend the deadline anymore

Comments

General Tips

- No need to understand the rationale behind the layers and model architecture. Just focus on the computation and data access patterns
- Whenever you modify something, make sure to check the correctness
 - It is very hard to debug when you make multiple parallelization/optimization at once
- Find out the problem **before** you do something
- Don't go directly into famous LLM inference optimization techniques... do the easy things first
- The skeleton code does many things inefficiently. Find them out and get rid of it

Things to Try (When you don't have any idea)

- Understand the skeleton code
- Parallelize the layers one-by-one, starting from the most time-consuming ones
 - There is an example CUDA kernel in the skeleton code. Check out `add_kernel1()` and `add_cuda()` in `src/layer.cu`
- Minimize memory allocation and free
- Minimize data movement
- Batch the inputs

Things to Try (When you think there's nothing to do)

- Profile your program
 - Try Nsight Systems profiler. We will cover this later in the class
 - Inspect the trace of your program. Find out any unexpected behavior
- Kernel optimization
 - How fast is your kernel compared to the peak FLOPS?
 - Try some kernel optimization techniques covered in the class
- Find out any redundant computations/data movement, and remove it

Updates - Code

[2024.05.09. 21:00] main.cpp 코드 수정 (eTL 공지)

[2024.05.12. 16:00] answer.bin 오류 수정, main.cpp 코드 수정 (fflush 추가)

[2024.06.10. 21:00] main.cpp 코드 수정 (validate 함수 수정 및 tolerance 값 상향 조정)

Updates - Restrictions

[2024.06.04. 12:00]

- ~~• Init phase에서 새로운 메모리 할당(malloc)은 가능하지만 메모리 복사(memcpy)는 불가능
 - ~~• E.g., input, output, activation 등을 미리 malloc 가능~~
 - ~~• E.g., input의 값을 다른 MPI rank로 memcpy (broadcast, scatter 등)는 불가능~~~~
- Batch size는 본인이 원하는 값으로 설정 가능 (레포트에 본인이 설정한 batch size 명시)
- Input prompt 개수(-n)은 임의로 선택 가능 (레포트에 본인이 선택한 n의 값 명시)
- FP16 사용 불가능 (Tensor Core 쓰지 말 것)

Updates - Restrictions (cont'd)

[2024.06.05. 14:00]

- (p.22 restriction 수정사항) 최종 결과가 같다면 불필요한 연산은 생략 가능
 - E.g., LM head의 linear 연산 중 일부
- (p.29 restriction 수정사항) Init phase에서 알 수 있는 정보로 할 수 있는 작업은 모두 허용